

Practical No.3

Data Science and Visualization (Honors Course)

Name: Preeti Shantilal Chaudhari

PRN: 72017986G

Class: TE ENT C 'B'

To determine the age of abalone on the basis of its physical measurements

In [19]:

```
import pandas as pd
```

In [20]:

```
col = ['sex', 'length', 'diameter', 'height', 'weight', 'sweight', 'vweight', 'shweight', 'rings']  
df=pd.read_csv('abalone.csv')
```

In [21]:

```
df.head()
```

Out[21]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 15 |
| 1 | M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 7 |
| 2 | F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 9 |
| 3 | M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 10 |
| 4 | I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 7 |

In [22]:

```
df.describe()
```

Out[22]:

| | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight |
|-------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 |
| mean | 0.523992 | 0.407881 | 0.139516 | 0.828742 | 0.359367 | 0.180594 | 0.238541 |
| std | 0.120093 | 0.099240 | 0.041827 | 0.490389 | 0.221963 | 0.109614 | 0.139501 |
| min | 0.075000 | 0.055000 | 0.000000 | 0.002000 | 0.001000 | 0.000500 | 0.001000 |
| 25% | 0.450000 | 0.350000 | 0.115000 | 0.441500 | 0.186000 | 0.093500 | 0.130000 |
| 50% | 0.545000 | 0.425000 | 0.140000 | 0.799500 | 0.336000 | 0.171000 | 0.234000 |
| 75% | 0.615000 | 0.480000 | 0.165000 | 1.153000 | 0.502000 | 0.253000 | 0.329000 |
| max | 0.815000 | 0.650000 | 1.130000 | 2.825500 | 1.488000 | 0.760000 | 1.005000 |

We can say the dataset here is already cleaned because there are no null values.

In [24]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Sex              4177 non-null   object
1   Length           4177 non-null   float64
2   Diameter         4177 non-null   float64
3   Height           4177 non-null   float64
4   Whole weight     4177 non-null   float64
5   Shucked weight   4177 non-null   float64
6   Viscera weight   4177 non-null   float64
7   Shell weight     4177 non-null   float64
8   Rings            4177 non-null   int64
dtypes: float64(7), int64(1), object(1)
memory usage: 293.8+ KB
```

In [12]:

```
X = df.drop('Rings' , axis=1) #Input
y = df['Rings'] #Output
```

In [13]:

```
X.head()
```

Out[13]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|
| 0 | M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 |
| 1 | M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 |
| 2 | F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 |
| 3 | M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 |
| 4 | I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 |

In [14]:

```
from collections import Counter  
Counter(y)
```

Out[14]:

```
Counter({15: 103,  
        7: 391,  
        9: 689,  
        10: 634,  
        8: 568,  
        20: 26,  
        16: 67,  
        19: 32,  
        14: 126,  
        11: 487,  
        12: 267,  
        18: 42,  
        13: 203,  
        5: 115,  
        4: 57,  
        6: 259,  
        21: 14,  
        17: 58,  
        22: 6,  
        1: 1,  
        3: 15,  
        26: 1,  
        23: 9,  
        29: 1,  
        2: 1,  
        27: 2,  
        25: 1,  
        24: 2})
```

In [17]:

```
set(X['Sex']) #Displaying unique entries
```

Out[17]:

```
{'F', 'I', 'M'}
```

In [26]:

```
from sklearn.preprocessing import LabelEncoder
enc=LabelEncoder()
X['Sex']=enc.fit_transform(X['Sex'])
```

In [27]:

```
set(X['Sex'])
```

Out[27]:

```
{0, 1, 2}
```

In [28]:

```
df.head()
```

Out[28]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|---|-----|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0 | M | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 15 |
| 1 | M | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 7 |
| 2 | F | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 9 |
| 3 | M | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 10 |
| 4 | I | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 7 |

In [29]:

```
from sklearn.model_selection import train_test_split
```

In [34]:

```
X_train,X_test,y_train,y_test = train_test_split(X,y,random_state=0,test_size=0.25)
#Splitting the dataset
```

In [33]:

```
len(X_train)
```

Out[33]:

```
3132
```

In [35]:

```
len(X_test)
```

Out[35]:

1045

In [36]:

```
X_train.head()
```

Out[36]:

| | Sex | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight |
|------|-----|--------|----------|--------|--------------|----------------|----------------|--------------|
| 940 | 1 | 0.460 | 0.345 | 0.105 | 0.4490 | 0.1960 | 0.0945 | 0.1265 |
| 2688 | 2 | 0.630 | 0.465 | 0.150 | 1.0270 | 0.5370 | 0.1880 | 0.1760 |
| 1948 | 2 | 0.635 | 0.515 | 0.165 | 1.2290 | 0.5055 | 0.2975 | 0.3535 |
| 713 | 2 | 0.355 | 0.265 | 0.085 | 0.2010 | 0.0690 | 0.0530 | 0.0695 |
| 3743 | 0 | 0.705 | 0.555 | 0.195 | 1.7525 | 0.7105 | 0.4215 | 0.5160 |

Prediction

In [37]:

```
from sklearn.naive_bayes import GaussianNB
```

In [38]:

```
clf = GaussianNB()
```

In [39]:

```
#train  
clf.fit(X_train,y_train)
```

Out[39]:

```
GaussianNB()
```

In [40]:

```
y_pred=clf.predict(X_test)
```

In [42]:

```
from sklearn.metrics import accuracy_score  
from sklearn.metrics import classification_report
```

In [43]:

```
accuracy_score(y_test,y_pred)*100
```

Out[43]:

26.02870813397129

The accuracy score is low due to presence of multiple classes.

Regression

$\text{precision} = \text{TP} / (\text{TP} + \text{FP})$

$\text{recall} = \text{TP} / (\text{TP} + \text{FN})$

$\text{f1-score} = 2\text{PR} / (\text{P} + \text{R})$

Support is the number of actual occurrences of class in a specified dataset.

In [44]:

```
print(classification_report(y_test,y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 3 | 0.50 | 1.00 | 0.67 | 7 |
| 4 | 0.30 | 0.62 | 0.40 | 13 |
| 5 | 0.27 | 0.42 | 0.33 | 40 |
| 6 | 0.32 | 0.43 | 0.36 | 63 |
| 7 | 0.26 | 0.36 | 0.30 | 114 |
| 8 | 0.27 | 0.29 | 0.28 | 139 |
| 9 | 0.25 | 0.30 | 0.27 | 152 |
| 10 | 0.21 | 0.24 | 0.23 | 139 |
| 11 | 0.26 | 0.42 | 0.32 | 121 |
| 12 | 0.50 | 0.01 | 0.02 | 93 |
| 13 | 0.00 | 0.00 | 0.00 | 51 |
| 14 | 0.00 | 0.00 | 0.00 | 32 |
| 15 | 0.00 | 0.00 | 0.00 | 22 |
| 16 | 0.00 | 0.00 | 0.00 | 16 |
| 17 | 0.00 | 0.00 | 0.00 | 12 |
| 18 | 0.00 | 0.00 | 0.00 | 6 |
| 19 | 0.00 | 0.00 | 0.00 | 10 |
| 20 | 0.00 | 0.00 | 0.00 | 8 |
| 21 | 0.00 | 0.00 | 0.00 | 2 |
| 22 | 0.00 | 0.00 | 0.00 | 1 |
| 23 | 0.00 | 0.00 | 0.00 | 2 |
| 24 | 0.00 | 0.00 | 0.00 | 1 |
| 27 | 0.00 | 0.00 | 0.00 | 0 |
| 29 | 0.00 | 0.00 | 0.00 | 1 |
| | | | | |
| accuracy | | | 0.26 | 1045 |
| macro avg | 0.13 | 0.17 | 0.13 | 1045 |
| weighted avg | 0.24 | 0.26 | 0.22 | 1045 |

C:\Users\HP\anaconda3\lib\site-packages\sklearn\metrics_classification.p
y:1245: UndefinedMetricWarning: Precision and F-score are ill-defined and
being set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\metrics_classification.p
y:1245: UndefinedMetricWarning: Recall and F-score are ill-defined and be
ing set to 0.0 in labels with no true samples. Use `zero_division` parame
ter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\metrics_classification.p
y:1245: UndefinedMetricWarning: Precision and F-score are ill-defined and
being set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\metrics_classification.p
y:1245: UndefinedMetricWarning: Recall and F-score are ill-defined and be
ing set to 0.0 in labels with no true samples. Use `zero_division` parame
ter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\HP\anaconda3\lib\site-packages\sklearn\metrics_classification.p
y:1245: UndefinedMetricWarning: Precision and F-score are ill-defined and
being set to 0.0 in labels with no predicted samples. Use `zero_division`
parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
C:\Users\HP\anaconda3\lib\site-packages\sklearn\metrics\_classification.p
y:1245: UndefinedMetricWarning: Recall and F-score are ill-defined and be
ing set to 0.0 in labels with no true samples. Use `zero_division` parame
ter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))
```

In [45]:

```
from sklearn.linear_model import LinearRegression
```

In [46]:

```
reg=LinearRegression()
```

In [47]:

```
reg.fit(X_train,y_train)
```

Out[47]:

```
LinearRegression()
```

In [48]:

```
y_pred = reg.predict(X_test)
```

In [49]:

```
y_pred
```

Out[49]:

```
array([13.10451425,  9.66747548, 10.35605247, ...,  9.95962005,
        12.59111443, 12.18516586])
```

In [50]:

```
from sklearn.metrics import mean_absolute_error
```

In [51]:

```
mean_absolute_error(y_test,y_pred) #summation of (|y_pred-y_train|/no.of entries)
```

Out[51]:

```
1.5955158378194019
```

In [52]:

```
from sklearn.metrics import r2_score
```


In [53]:

```
r2_score(y_test,y_pred)    #r2_score = 1-(summation of (y_pred-y_train)^2 / summation of (me
```

Out[53]:

0.5354158501894077

In this case we can say that Regression outperforms GaussianNB in terms of accuracy. *(due to the dataset)*

In []: